


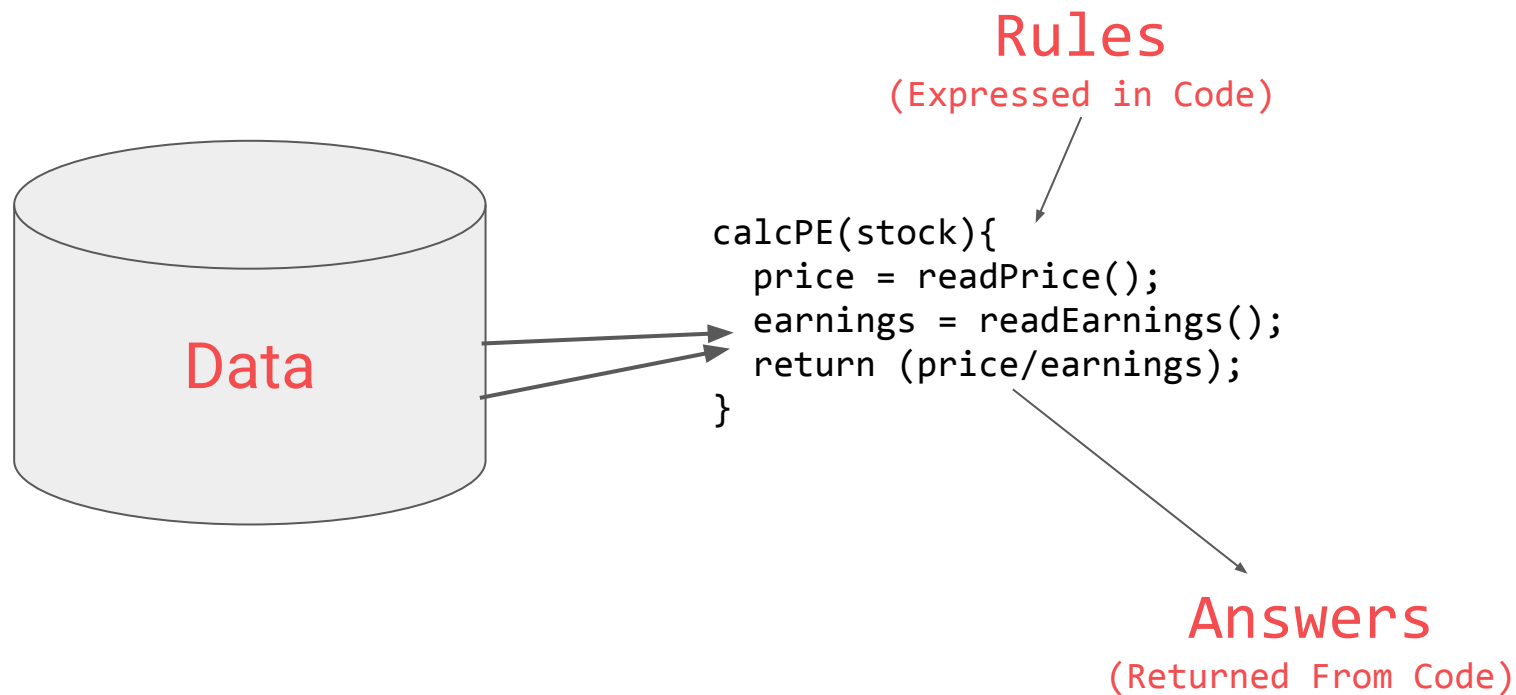



Zero to ~~Here~~ Hinton with TensorFlow

Arnaldo Gualberto & Mikaeri Ohana
Google Developer Experts in ML









```
if (ball.collide(brick)){  
    removeBrick();  
    ball.dx=-1*(ball.dx);  
    ball.dy=-1*(ball.dy);  
}
```







Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```





Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



```
// ????
```



Activity Recognition



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING



1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011

Label = RUNNING



10010100111111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING



11111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING
(Sort of)



Activity Recognition



```
0101001010100101010  
1001010101001011101  
0100101010010101001  
0101001010100101010
```

Label = WALKING



```
1010100101001010101  
0101010010010010001  
0010011111010101111  
1010100100111101011
```

Label = RUNNING



```
1001010011111010101  
1101010111010101110  
1010101111010101011  
1111110001111010101
```

Label = BIKING



```
1111111111010011101  
0011111010111110101  
0101110101010101110  
1010101010100111110
```

Label = GOLFING
(Sort of)



Activity Recognition



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING



1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011

Label = RUNNING



10010100111111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING



11111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING
(Sort of)




The Machine Learning Paradigm

Make a Guess!

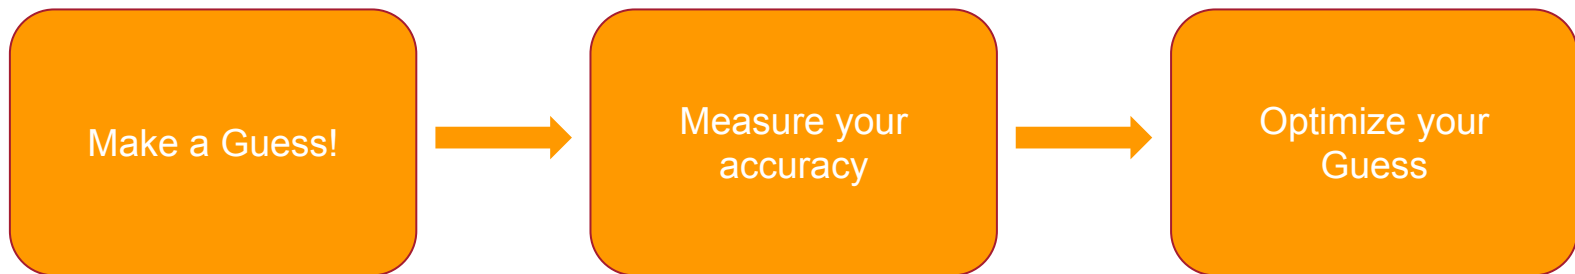


The Machine Learning Paradigm





The Machine Learning Paradigm






The Machine Learning Paradigm



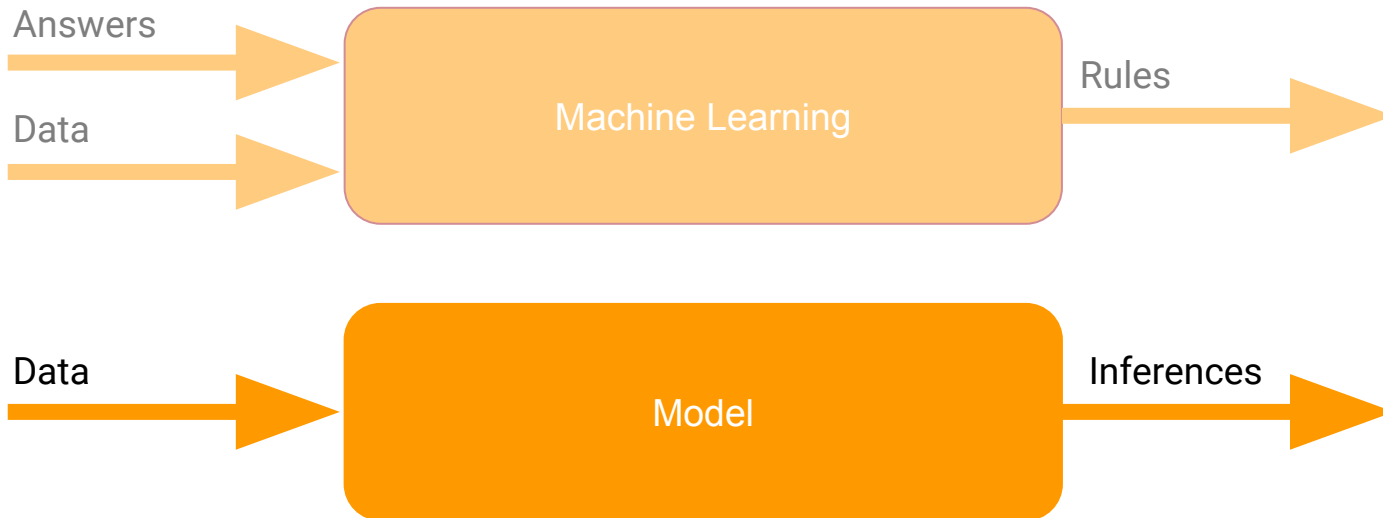


The Machine Learning Paradigm





The Machine Learning Paradigm





Computer Vision

Using Machine Learning to understand images



hit the button



**WHEN DID YOU BECOME AN EXPERT IN
COMPUTER VISION?**



LAST NIGHT




TODAY















+






+




+






+



+




=


HUMAN



+




+




=


HUMAN



+

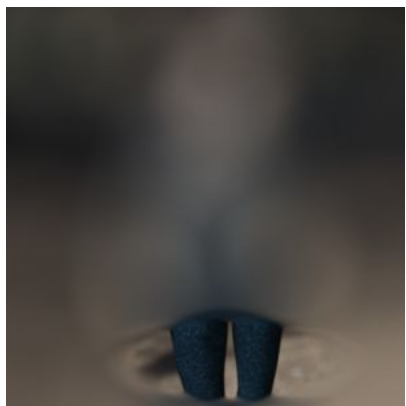


+




=


HORSE



+




+




=


HUMAN



+




+




=


HORSE



+




+




=

HUMAN




+




=

HORSE







Filters can then be combined with labels to make a prediction of the image contents...



Filters can then be combined with labels to make a prediction of the image contents...

The filters that match the label are learned over time!





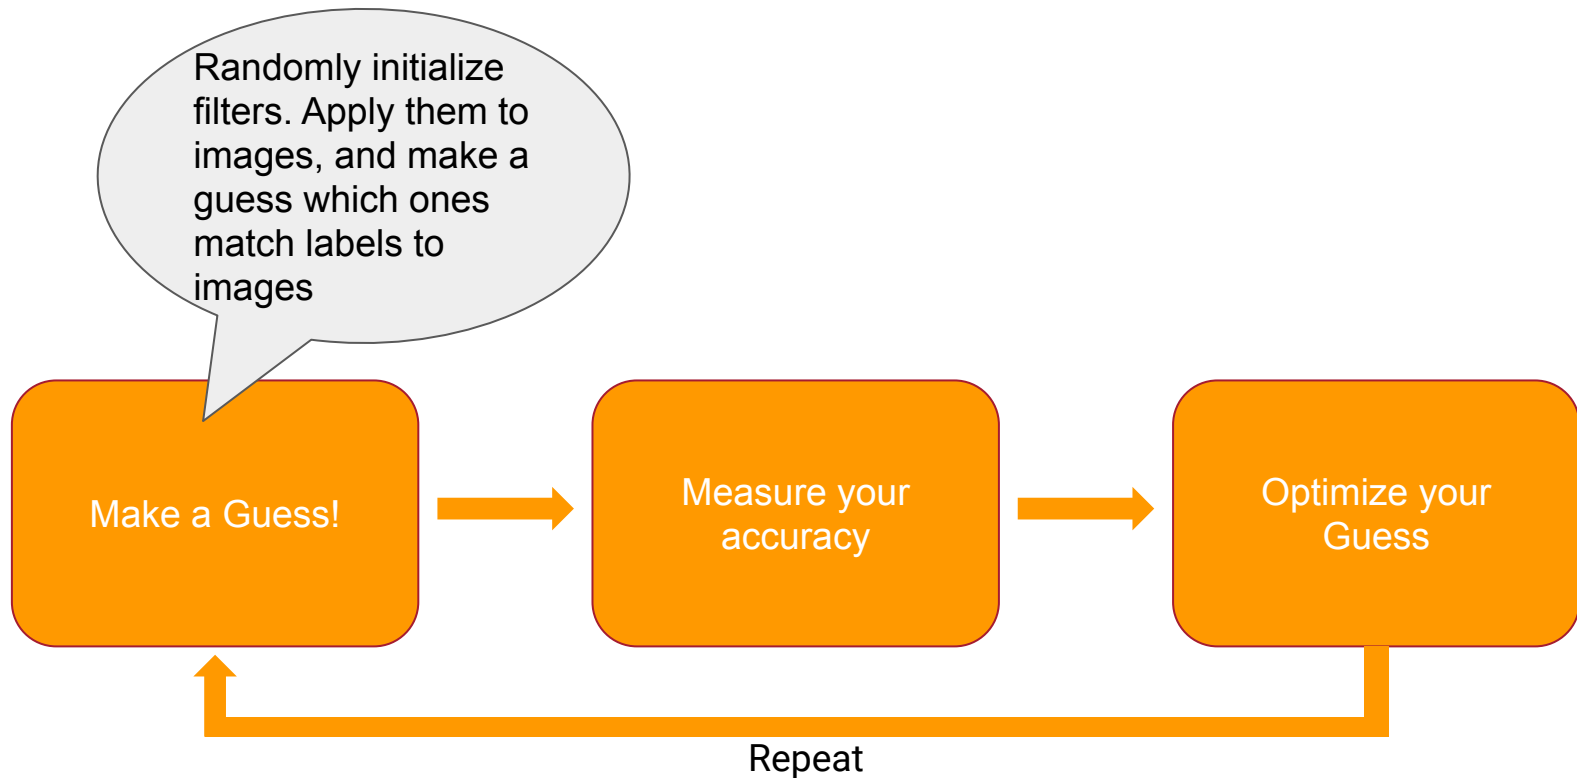


The Machine Learning Paradigm



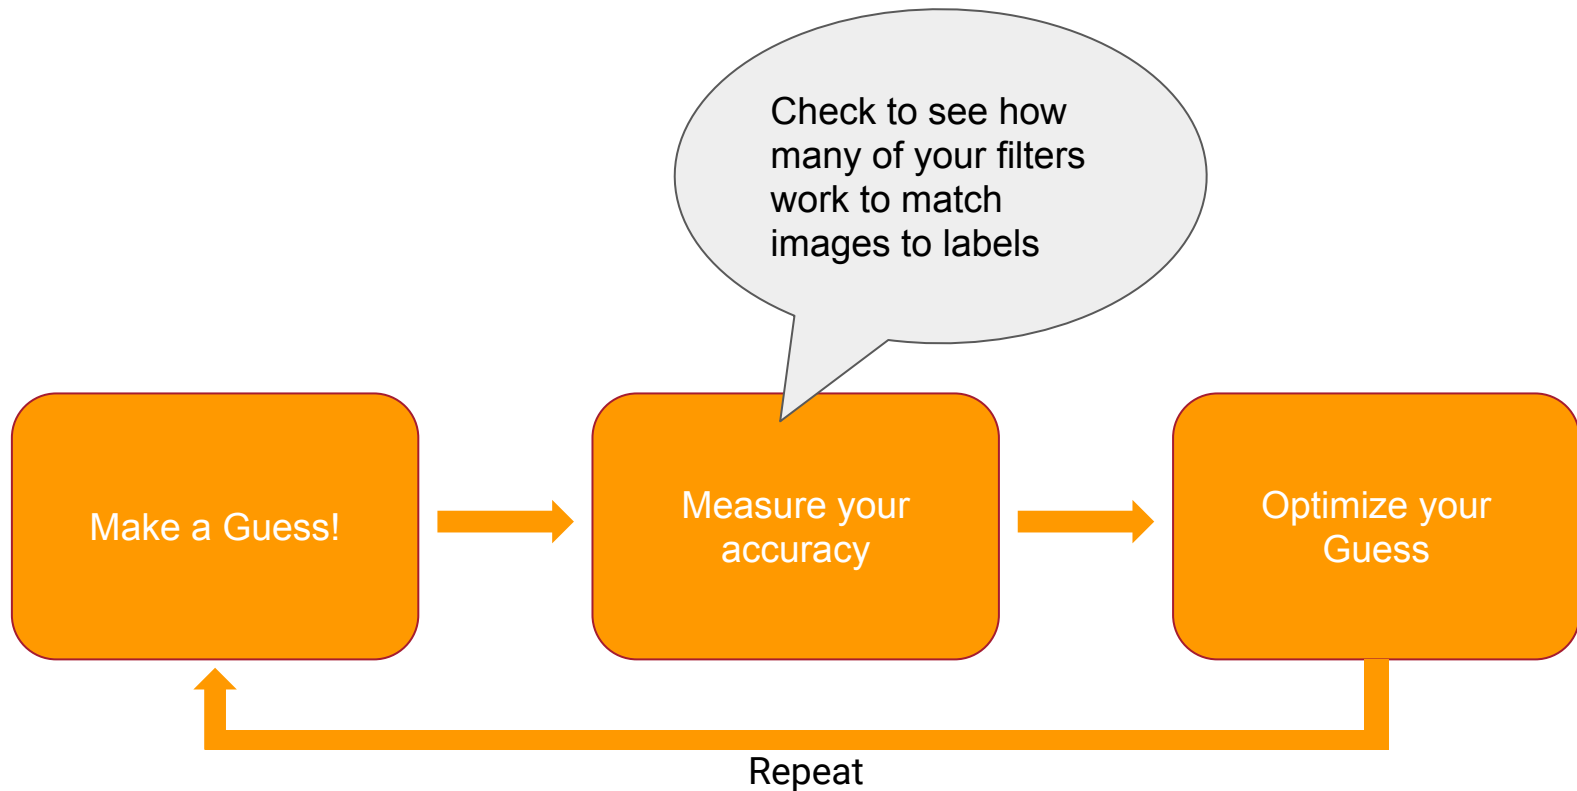


The Machine Learning Paradigm



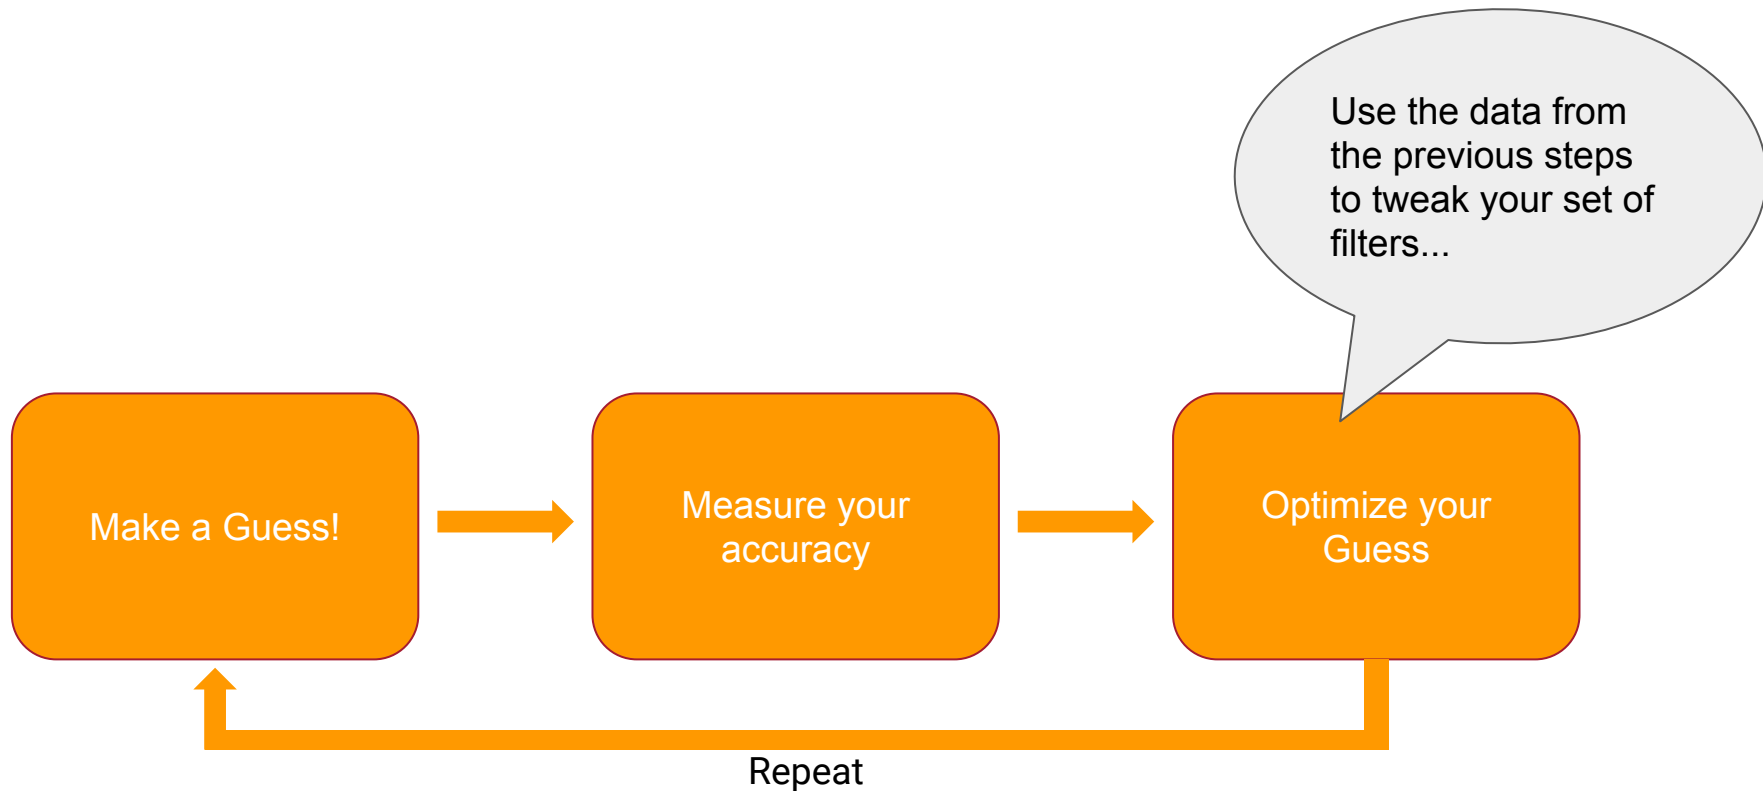


The Machine Learning Paradigm



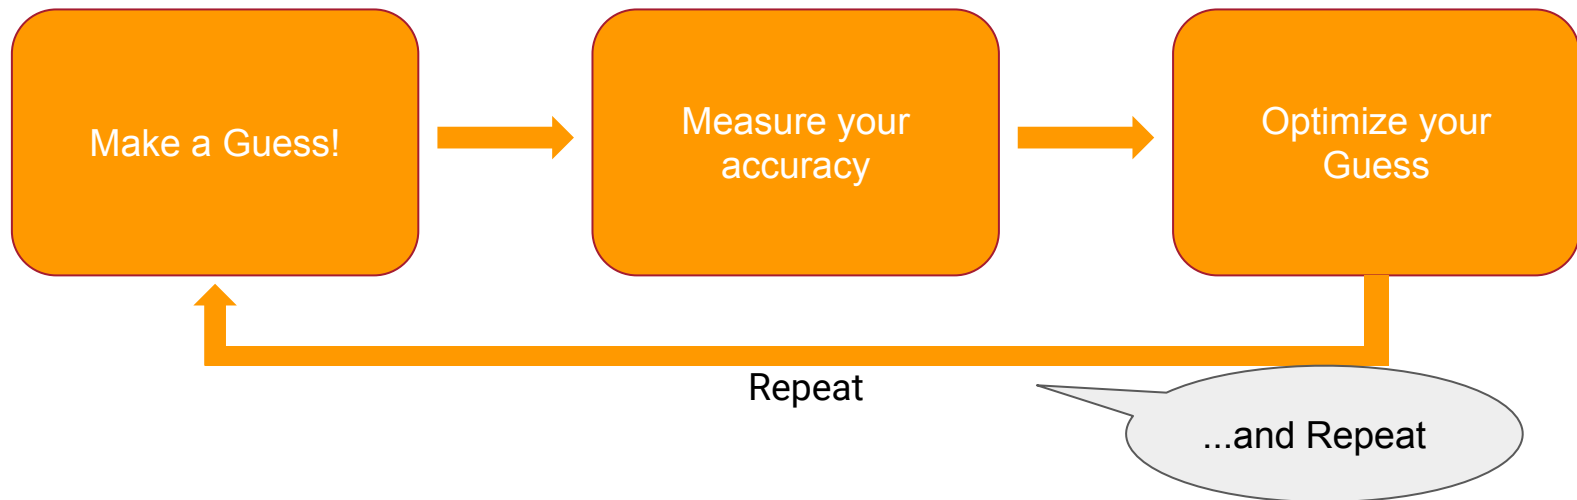




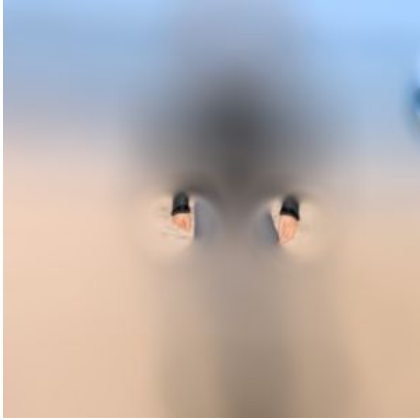
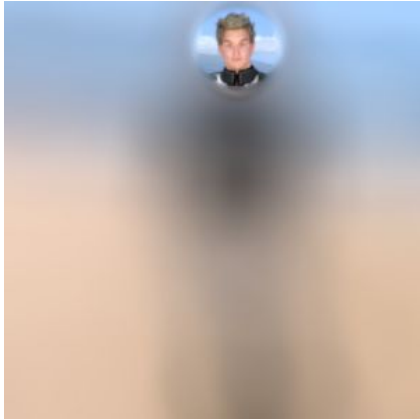

The Machine Learning Paradigm







The Machine Learning Paradigm




  +  +   = ?






+



+



=

HUMAN



Verizon

4:20 PM

76%

< Albums

chihuahua or muffin

Select



Verizon LTE

3:02 PM

42%

< Albums

sheepdog or mop


Select



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
        input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Conv2D stands
for 2D
Convolution --
another word for
a filter



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
        input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

MaxPooling is a way of compressing the image while enhancing features

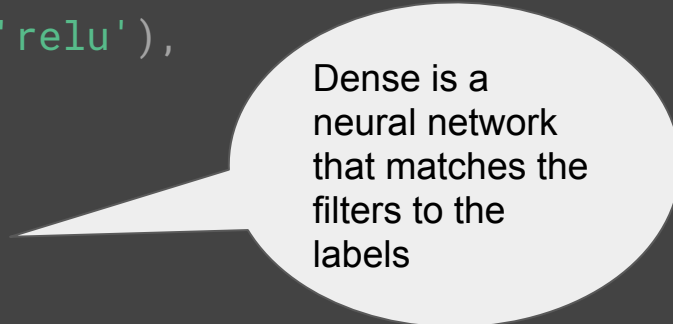
1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Feature map



Pooled
Feature map

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
        input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```





Dense is a
neural network
that matches the
filters to the
labels

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',  
        input_shape=(300, 300, 3)),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(512, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

The final Dense
represents the
labels:
Horse = 0,
Human = 1

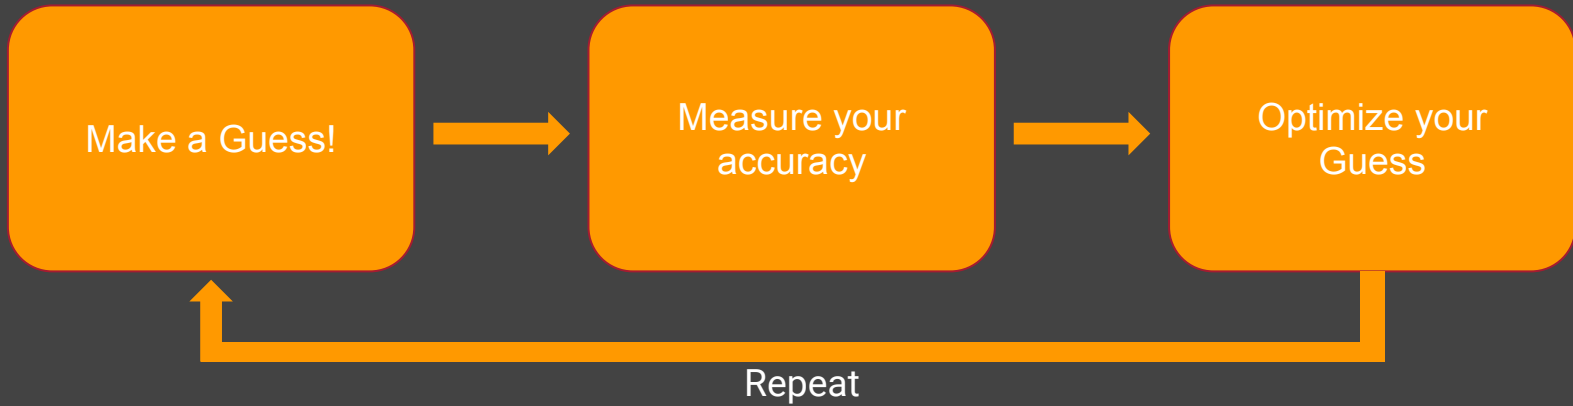

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
        input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

ReLU and Sigmoid
are non-linear
activation
functions!

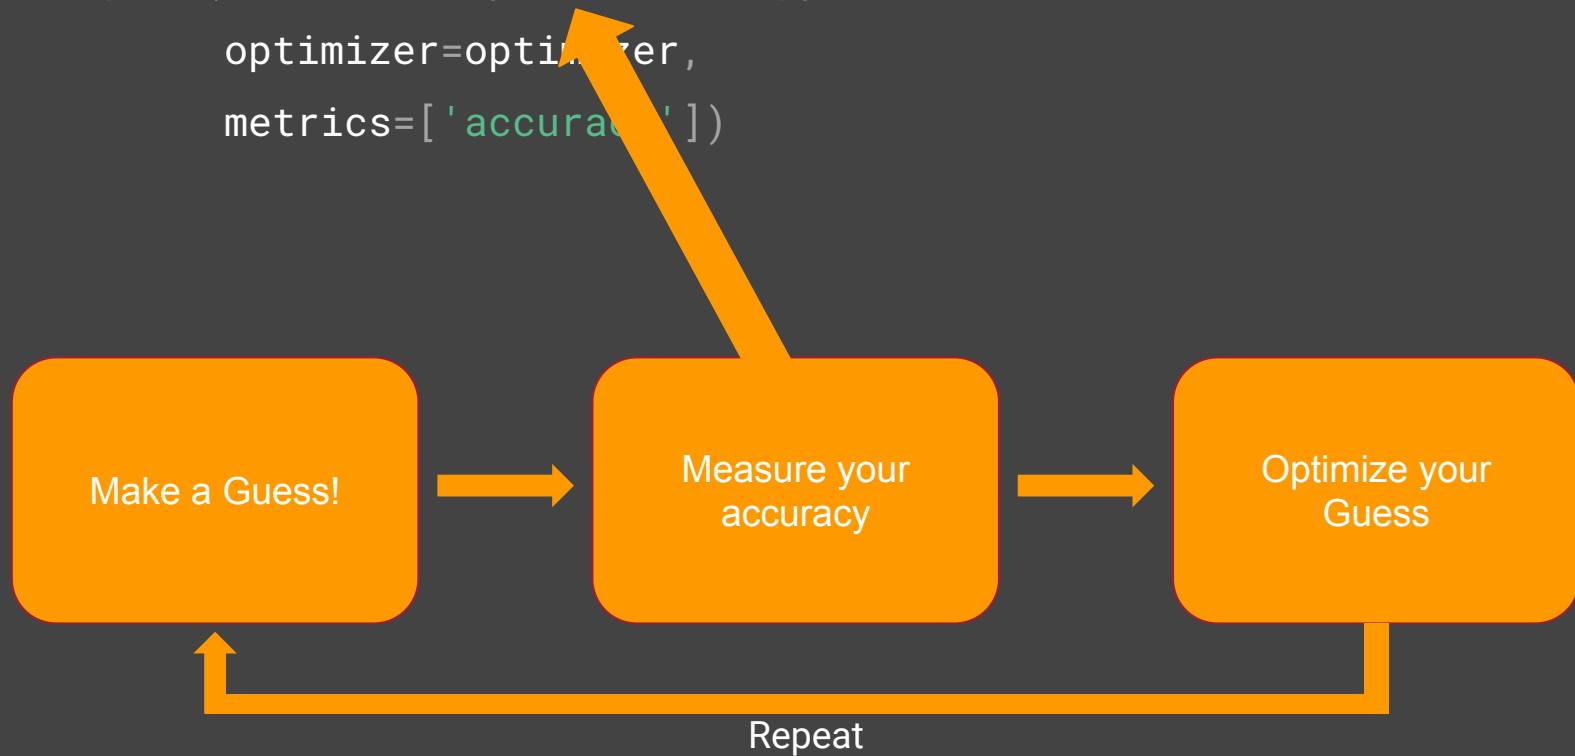


```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```


```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```



```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```

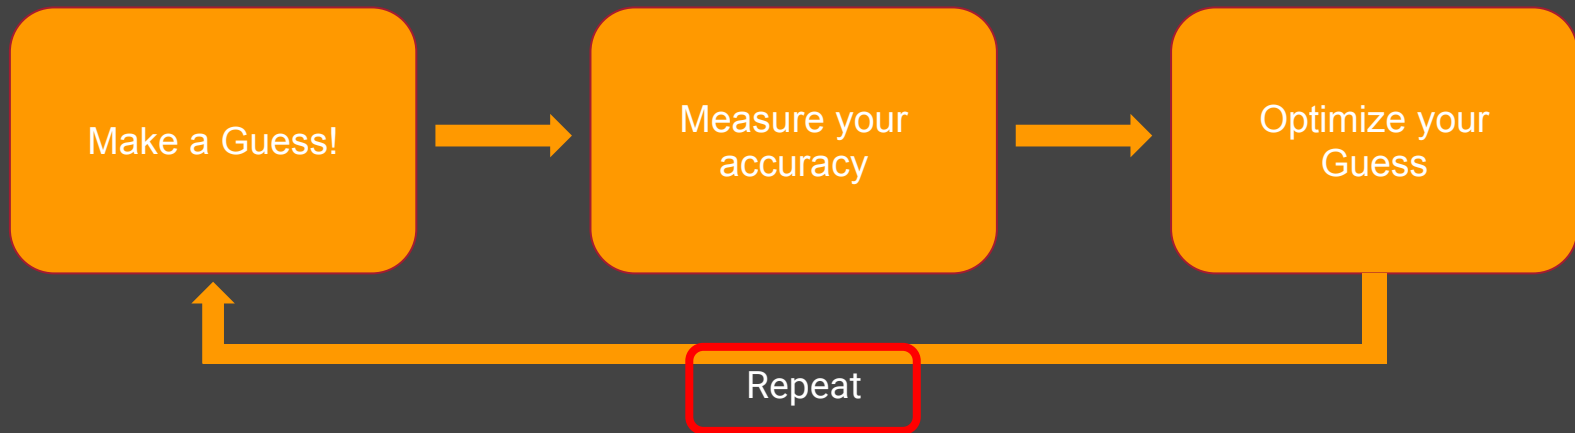


```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'],
```



```
model.fit(train_generator, epochs=12, ...)
```

```
model.fit(train_generator, epochs=12, ...)
```





Demo

Training a computer to recognize Horses or Humans



Thank you!



@miohana



@miohana



@mikaeriohana



@mikaeriohana



@arnaldo.g12



@arnaldog12



@arnaldog12_



@arnaldog12